

Knowledge Representation & Reasoning



Knowledge Representation

Dr Ruairi O'Reilly

To be covered:

- A general ontology, which organizes everything in the world into a hierarchy of categories.
- The basic categories of objects, substances, and measures.
- Events.
- Knowledge about beliefs.
- Reasoning systems designed for efficient inference with categories.
- Reasoning with default information. (TBC)
- Finally we bring all the knowledge together in the context of an Internet shopping environment. (TBC)

Ontological Engineering

Toy domains vrs Complex domains?

Examples of Complex domains (shopping on the Internet or driving a car in traffic)

=> more general and flexible representations.

How to create these representations, concentrating on general concepts—such as Events, Time, Physical Objects, and Beliefs— that occur in many different domains. Representing these abstract concepts is sometimes called ontological engineering.

Representing everything in the world....

Upper Ontology

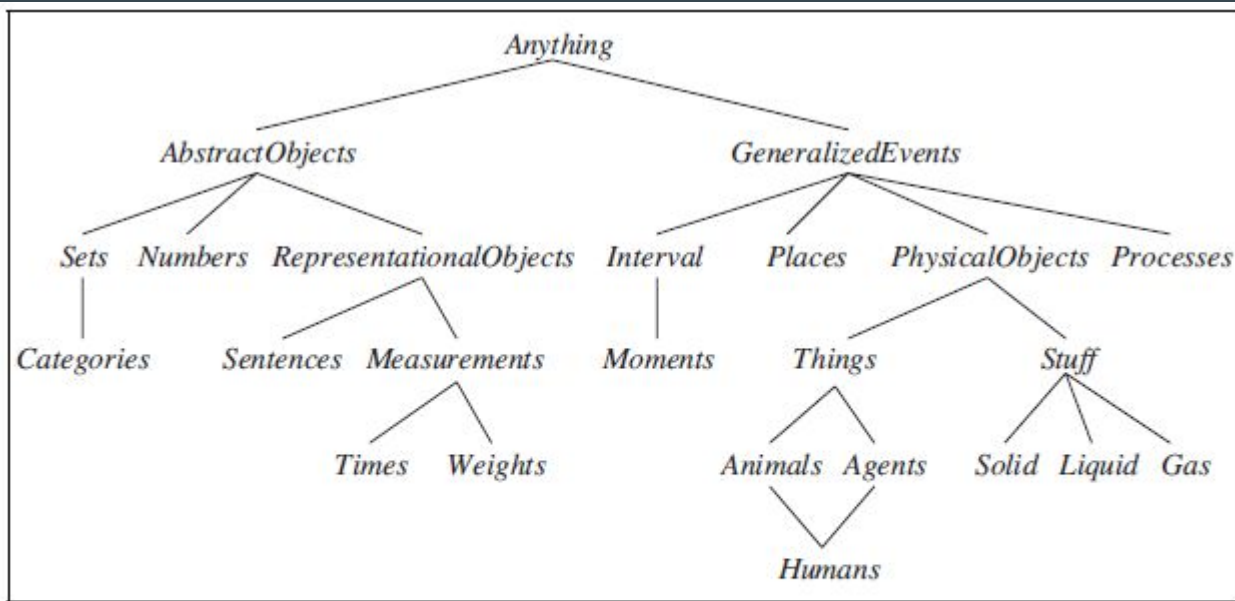


Figure 12.1 The upper ontology of the world, showing the topics to be covered later in the chapter. Each link indicates that the lower concept is a specialization of the upper one. Specializations are not necessarily disjoint; a human is both an animal and an agent, for example. We will see in Section 12.3.3 why physical objects come under generalized events.

One important
caveat...

Of what use is an upper ontology? Wumpus World

Although we do represent time, it has a simple structure: Nothing happens except when the agent acts, and all changes are instantaneous.

A more general ontology, better suited for the real world, would allow for simultaneous changes extended over time. We also used a Pit predicate to say which squares have pits. We could have allowed for different kinds of pits by having several individuals belonging to the class of pits, each having different properties. Similarly, we might want to allow for other animals besides wumpuses. It might not be possible to pin down the exact species from the available percepts, so we would need to build up a biological taxonomy to help the agent predict the behavior of cave-dwellers from scanty clues.

General Purpose Ontology

Two major characteristics of general-purpose ontologies distinguish them from collections of special-purpose ontologies:

- A general-purpose ontology should be applicable in more or less any special-purpose domain (with the addition of domain-specific axioms). This means that no representational issue can be finessed or brushed under the carpet.
- In any sufficiently demanding domain, different areas of knowledge must be unified, because reasoning and problem solving could involve several areas simultaneously.

None of the top AI applications make use of a shared ontology—they all use special-purpose knowledge engineering. Social/political considerations can make it difficult for competing parties to agree on an ontology.

Categories And Objects

Choices for representing categories in first-order logic: predicates and objects

Basketball (b), or we can reify the category as an object, Basketballs. We could then say $\text{Member}(b, \text{Basketballs})$, which we will abbreviate as $b \in \text{Basketballs}$, to say that b is a member of the category of basketballs. We say $\text{Subset}(\text{Basketballs}, \text{Balls})$, abbreviated as $\text{Basketballs} \subset \text{Balls}$, to say that Basketballs is a subcategory of Balls. We will use subcategory, subclass, and subset interchangeably.

Categories serve to organize and simplify the knowledge base through inheritance.

Subclass relations organize categories into a taxonomy, or taxonomic hierarchy. Taxonomies have been used explicitly for centuries in technical fields.

First-order logic makes it easy to state facts about categories, either by relating objects to categories or by quantifying over their members. Here are some types of facts, with examples of each:

- An object is a member of a category.

$BB9 \in \text{Basketballs}$

- A category is a subclass of another category.

$\text{Basketballs} \subset \text{Balls}$

- All members of a category have some properties.

$(x \in \text{Basketballs}) \Rightarrow \text{Spherical}(x)$

- Members of a category can be recognized by some properties.

$\text{Orange}(x) \wedge \text{Round}(x) \wedge \text{Diameter}(x)=9.5 \wedge x \in \text{Balls} \Rightarrow x \in \text{Basketballs}$

- A category as a whole has some properties.

$\text{Dogs} \in \text{DomesticatedSpecies}$

State relations between categories that are not subclasses

For example, if we just say that Males and Females are subclasses of Animals, then we have not said that a male cannot be a female. We say that two or more categories are disjoint if they have no members in common. And even if we know that males and females are disjoint, we will not know that an animal that is not a male must be a female, unless we say that males and females constitute an exhaustive decomposition of the animals. A disjoint exhaustive decomposition is known as a partition. The following examples illustrate these three concepts:

Disjoint({Animals, Vegetables})

ExhaustiveDecomposition({Americans, Canadians, Mexicans}, NorthAmericans)

Partition({Males, Females}, Animals)

(Note that the ExhaustiveDecomposition of NorthAmericans is not a Partition, because some people have dual citizenship.)

Physical composition

The idea that one object can be part of another is a familiar one. Objects can be grouped into PartOf hierarchies, reminiscent of the Subset hierarchy:

i) *PartOf (Bucharest , Romania)*

ii) *PartOf (Romania, EasternEurope)*

iii) *PartOf (EasternEurope, Europe)*

iv) *PartOf (Europe, Earth) .*

The PartOf relation is transitive and reflexive; that is,

$PartOf(x, y) \wedge PartOf(y, z) \Rightarrow PartOf(x, z) .$

$PartOf(x, x) .$

Measurements

In both scientific and commonsense theories of the world, objects have height, mass, cost, and so on. The values that we assign for these properties are called measures. Ordinary quantitative measures are quite easy to represent. We imagine that the universe includes abstract “measure objects,” such as the length that is the length of a 1.5 inch line segment: For instance we can say the length is 1.5 inches or 3.81 centimeters. Thus, the same length has different names in our language. We represent the length with a units function that takes a number as argument.

$$\textit{Length}(L1)=\textit{Inches}(1.5)=\textit{Centimeters}(3.81)$$

Conversion between units is done by equating multiples of one unit to another:

$$\textit{Centimeters}(2.54 \times d)=\textit{Inches}(d) .$$

Similar axioms can be written for pounds and kilograms, seconds and days, and dollars and cents. Measures can be used to describe objects as follows:

Diameter (Basketball 12)=Inches(9.5) .

ListPrice(Basketball 12)=\$(19) .

$d \in \text{Days} \Rightarrow \text{Duration}(d)=\text{Hours}(24)$

Note that \$(1) is not a dollar bill! One can have two dollar bills, but there is only one object named \$(1). Note also that, while Inches(0) and Centimeters(0) refer to the same zero length, they are not identical to other zero measures, such as Seconds(0).

Objects: Things and stuff

The real world can be seen as consisting of primitive objects (e.g., atomic particles) and composite objects built from them. By reasoning at the level of large objects such as apples and cars, we can overcome the complexity involved in dealing with vast numbers of primitive objects individually. There is, however, a significant portion of reality that seems to defy any obvious individuation—division into distinct objects. We give this portion the generic name stuff. For example, suppose I have some butter and an aardvark in front of me. I can say there is one aardvark, but there is no obvious number of “butter-objects,” because any part of a butter-object is also a butter-object, at least until we get to very small parts indeed. This is the major distinction between stuff and things. If we cut an aardvark in half, we do not get two aardvarks (unfortunately).

Count nouns vrs Mass nouns

With some caveats about very small parts that we will omit for now, any part of a butter-object is also a butter-object:

$$b \in \textit{Butter} \wedge \textit{PartOf}(p, b) \Rightarrow p \in \textit{Butter} .$$

We can now say that butter melts at around 30 degrees centigrade:

$$b \in \textit{Butter} \Rightarrow \textit{MeltingPoint}(b, \textit{Centigrade}(30)) .$$

Events

Situation calculus is limited in its applicability: it was designed to describe a world in which actions are discrete, instantaneous, and happen one at a time. Consider a continuous action, such as filling a bathtub. Situation calculus can say that the tub is empty before the action and full when the action is done, but it can't talk about what happens during the action. It also can't describe two actions happening at the same time—such as brushing one's teeth while waiting for the tub to fill. To handle such cases we introduce an alternative formalism known as event calculus, which is based on points of time rather than on situations.

Event calculus reifies fluents and events. The fluent $\text{At}(\text{Shankar}, \text{Berkeley})$ is an object that refers to the fact of Shankar being in Berkeley, but does not by itself say anything about whether it is true. To assert that a fluent is actually true at some point in time we use the predicate T , as in $T(\text{At}(\text{Shankar}, \text{Berkeley}), t)$.

Events are described as instances of event categories. The event $E1$ of Shankar flying from San Francisco to Washington, D.C. is described as

$$E1 \in \text{Flyings} \wedge \text{Flyer}(E1, \text{Shankar}) \wedge \text{Origin}(E1, \text{SF}) \wedge \text{Destination}(E1, \text{DC}) .$$

If this is too verbose, we can define an alternative three-argument version of the category of flying events and say

$$E1 \in \text{Flyings}(\text{Shankar}, \text{SF}, \text{DC})$$

We then use $\text{Happens}(E1, i)$ to say that the event $E1$ took place over the time interval i , and we say the same thing in functional form with $\text{Extent}(E1)=i$. We represent time intervals by a $(\text{start}, \text{end})$ pair of times; that is, $i = (t1, t2)$ is the time interval that starts at $t1$ and ends at $t2$.

To complete set of predicates for one version of the event calculus is

$T(f, t)$ Fluent f is true at time t

$Happens(e, i)$ Event e happens over the time interval i

$Initiates(e, f, t)$ Event e causes fluent f to start to hold at time t

$Terminates(e, f, t)$ Event e causes fluent f to cease to hold at time t

$Clipped(f, i)$ Fluent f ceases to be true at some point during time interval i

$Restored(f, i)$ Fluent f becomes true sometime during time interval i

We assume a distinguished event, $Start$, that describes the initial state by saying which fluents are initiated or terminated at the start time. We define T by saying that a fluent holds at a point in time if the fluent was initiated by an event at some time in the past and was not made false (clipped) by an intervening event. A fluent does not hold if it was terminated by an event and not made true (restored) by another event.

Processes

The events we have seen so far are what we call discrete events— they have a definite structure. Shankar’s trip has a beginning, middle, and end. If interrupted halfway, the event would be something different—it would not be a trip from San Francisco to Washington, but instead a trip from San Francisco to somewhere over Kansas. On the other hand, the category of events denoted by Flyings has a different quality. If we take a small interval of Shankar’s flight, say, the third 20-minute segment (while he waits anxiously for a bag of peanuts), that event is still a member of Flyings. In fact, this is true for any subinterval.

Categories of events with this property are called process categories or liquid event categories. Any process e that happens over an interval also happens over any subinterval: $(e \in \text{Processes}) \wedge \text{Happens}(e, (t_1, t_4)) \wedge (t_1 < t_2 < t_3 < t_4) \Rightarrow \text{Happens}(e, (t_2, t_3))$.

Time intervals

Event calculus opens us up to the possibility of talking about time, and time intervals. We will consider two kinds of time intervals: moments and extended intervals. The distinction is that only moments have zero duration:

Partition({Moments,ExtendedIntervals}, Intervals)

$i \in \text{Moments} \Leftrightarrow \text{Duration}(i) = \text{Seconds}(0)$

The function Duration gives the difference between the end time and the start time.

$\text{Interval}(i) \Rightarrow \text{Duration}(i) = (\text{Time}(\text{End}(i)) - \text{Time}(\text{Begin}(i)))$.

$\text{Time}(\text{Begin}(\text{AD1900})) = \text{Seconds}(0)$.

$\text{Time}(\text{Begin}(\text{AD2001})) = \text{Seconds}(3187324800)$.

$\text{Time}(\text{End}(\text{AD2001})) = \text{Seconds}(3218860800)$.

$\text{Duration}(\text{AD2001}) = \text{Seconds}(31536000)$.

To make these numbers easier to read, we also introduce a function *Date*, which takes six arguments (hours, minutes, seconds, day, month, and year) and returns a time point:

Time(Begin(AD2001))=Date(0, 0, 0, 1, Jan, 2001)

Date(0, 20, 21, 24, 1, 1995)=Seconds(3000000000) .

Two intervals Meet if the end time of the first equals the start time of the second. The complete set of interval relations, as proposed by Allen (1983) below:

Meet(i, j) ⇔ End(i)=Begin(j)

Before(i, j) ⇔ End(i) < Begin(j)

After (j, i) ⇔ Before(i, j)

During(i, j) ⇔ Begin(j) < Begin(i) < End(i) < End(j)

Overlap(i, j) ⇔ Begin(i) < Begin(j) < End(i) < End(j)

Begins(i, j) ⇔ Begin(i) = Begin(j)

Finishes(i, j) ⇔ End(i) = End(j)

Equals(i, j) ⇔ Begin(i) = Begin(j) ∧ End(i) = End(j)

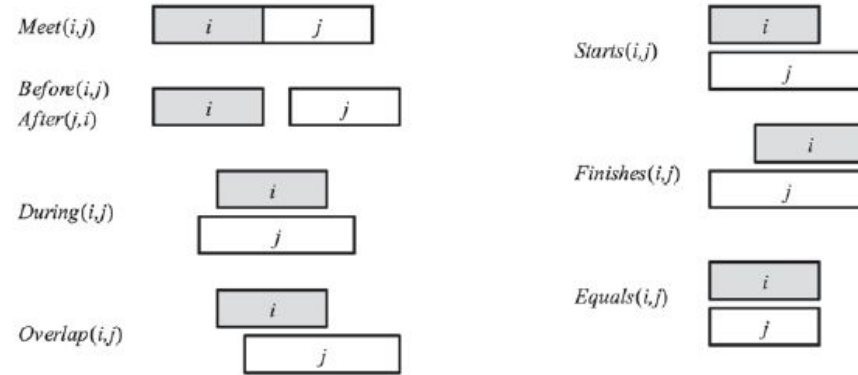


Figure 12.2 Predicates on time intervals.




Figure 12.3 A schematic view of the object *President(USA)* for the first 15 years of its existence.

Fluents & Objects

Physical objects can be viewed as generalized events, in the sense that a physical object is a chunk of space–time. For example, USA can be thought of as an event that began in, say, 1776 as a union of 13 states and is still in progress today as a union of 50. We can describe the changing properties of USA using state fluents, such as `Population(USA)`. A property of the USA that changes every four or eight years, barring mishaps, is its president. One might propose that `President(USA)` is a logical term that denotes a different object at different times. Unfortunately, this is not possible, because a term denotes exactly one object in a given model structure. (The term `President(USA, t)` can denote different objects, depending on the value of `t`, but our ontology keeps time indices separate from fluents.) The only possibility is that `President(USA)` denotes a single object that consists of different people at different times. It is the object that is George Washington from 1789 to 1797, John Adams from 1797 to 1801, and so on, as in Figure 12.3. To say that George Washington was president throughout 1790, we can write

Fluents & Objects

T(Equals(President(USA),GeorgeWashington),AD1790)

We use the function symbol Equals rather than the standard logical predicate =, because we cannot have a predicate as an argument to T, and because the interpretation is not that GeorgeWashington and President(USA) are logically identical in 1790; logical identity is not something that can change over time. The identity is between the subevents of each object that are defined by the period 1790.

That'll do folks....

Mental Events and Mental Objects

The agents we have constructed so far have beliefs and can deduce new beliefs. Yet none of them has any knowledge about beliefs or about deduction. Knowledge about one's own knowledge and reasoning processes is useful for controlling inference. For example, suppose Alice asks “what is the square root of 1764” and Bob replies “I don't know.” If Alice insists “think harder,” Bob should realize that with some more thought, this question can in fact be answered. On the other hand, if the question were “Is your mother sitting down right now?” then Bob should realize that thinking harder is unlikely to help. Knowledge about the knowledge of other agents is also important; Bob should realize that his mother knows whether she is sitting or not, and that asking her would be a way to find out.

We begin with the propositional attitudes that an agent can have toward mental objects: attitudes such as *Believes*, *Knows*, *Wants*, *Intends*, and *Informs*. The difficulty is that these attitudes do not behave like “normal” predicates. For example, suppose we try to assert that Lois knows that Superman can fly:

Knows(Lois, CanFly(Superman)) .

One minor issue with this is that we normally think of *CanFly(Superman)* as a sentence, but here it appears as a term. That issue can be patched up just by reifying

CanFly(Superman);

making it a fluent. A more serious problem is that, if it is true that Superman is Clark Kent, then we must conclude that Lois knows that Clark can fly:

(Superman = Clark) \wedge Knows(Lois , CanFly(Superman))

$\not\models$ Knows(Lois, CanFly(Clark)) .

Reasoning Systems for Categories

Categories are the primary building blocks of large-scale knowledge representation schemes. This section describes systems specially designed for organizing and reasoning with categories. There are two closely related families of systems: semantic networks provide graphical aids for visualizing a knowledge base and efficient algorithms for inferring properties of an object on the basis of its category membership; and description logics provide a formal language for constructing and combining category definitions and efficient algorithms for deciding subset and superset relationships between categories.

Semantic networks

There are many variants of semantic networks, but all are capable of representing individual objects, categories of objects, and relations among objects. A typical graphical notation displays object or category names in ovals or boxes, and connects them with labeled links. For example, Figure 12.5 has a *MemberOf* link between *Mary* and *FemalePersons*, corresponding to the logical assertion $Mary \in FemalePersons$; similarly, the *SisterOf* link between Mary and John corresponds to the assertion *SisterOf* (*Mary*, *John*). We can connect categories using *SubsetOf* links, and so on. It is such fun drawing bubbles and arrows that one can get carried away. For example, we know that persons have female persons as mothers, so can we draw a *HasMother* link from *Persons* to *FemalePersons*? The answer is no, because *HasMother* is a relation between a person and his or her mother, and categories do not have mothers.

Semantic networks

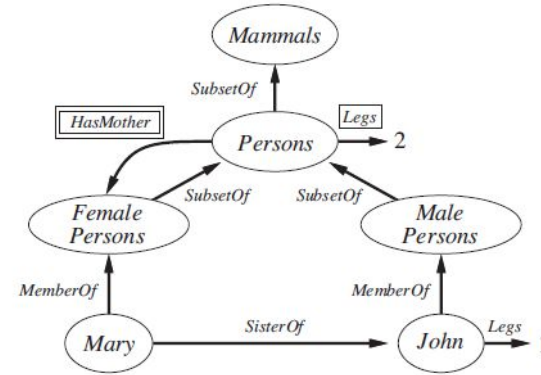


Figure 12.5 A semantic network with four objects (John, Mary, 1, and 2) and four categories. Relations are denoted by labeled links.




Figure 12.6 A fragment of a semantic network showing the representation of the logical assertion *Fly*(*Shankar*, *NewYork*, *NewDelhi*, *Yesterday*).

Description logics

The syntax of first-order logic is designed to make it easy to say things about objects. Description logics are notations that are designed to make it easier to describe definitions and properties of categories. Description logic systems evolved from semantic networks in response to pressure to formalize what the networks mean while retaining the emphasis on taxonomic structure as an organizing principle.

The principal inference tasks for description logics are subsumption (checking if one category is a subset of another by comparing their definitions) and classification (checking whether an object belongs to a category). Some systems also include consistency of a category definition—whether the membership criteria are logically satisfiable.

Description logics

Concept \rightarrow **Thing** | *ConceptName*
| **And**(*Concept*, ...)
| **All**(*RoleName*, *Concept*)
| **AtLeast**(*Integer*, *RoleName*)
| **AtMost**(*Integer*, *RoleName*)
| **Fills**(*RoleName*, *IndividualName*, ...)
| **SameAs**(*Path*, *Path*)
| **OneOf**(*IndividualName*, ...)
Path \rightarrow [*RoleName*, ...]

Figure 12.7 The syntax of descriptions in a subset of the CLASSIC language.

Reasoning with Default Information

See textbook for details.

The Internet Shopping World

See textbook for details.

References

All content is taken verbatim from Chapter 12 of AI: An Intelligent Approach 3rd edition

The three predicates are defined as follows:

$Disjoint(s) \Leftrightarrow (\forall c1, c2 \ c1 \in s \wedge c2 \in s \wedge c1 = c2 \Rightarrow Intersection(c1, c2) = \{ \})$

$ExhaustiveDecomposition(s, c) \Leftrightarrow (\forall i \ i \in c \Leftrightarrow \exists c2 \ c2 \in s \wedge i \in c2)$

$Partition(s, c) \Leftrightarrow Disjoint(s) \wedge ExhaustiveDecomposition(s, c)$

Categories can also be defined by providing necessary and sufficient conditions for membership. For example, a bachelor is an unmarried adult male:

$x \in Bachelors \Leftrightarrow Unmarried(x) \wedge x \in Adults \wedge x \in Males$

As we discuss in the sidebar on natural kinds on page 443, strict logical definitions for categories are neither always possible nor always necessary.